

P2L3 PThreads

PThreads == POSIX Threads

POSIX = Portable Operating System Interface

1. PThread Creation

```
1 // Correspond to Birrell's Thread datatype
2 pthread_t aThread; // Pthread datatype
3
4 // Correspond to Birrell's Fork(proc, args) for thread creation
5
6 // pthread_attr_t is a data structure that holds information about -
7 // certain things the thread manager should take into consideration -
8 // when schedule the thread
9
10 // returns status representation about whether the creation was success or failure.
11 int pthread_create (pthread_t *thread, const pthread_attr_t *attr,
12 void * (*start_routine)(void *), void *arg)
13
14 // Correspond to Join(thread)
15
16 // status captures all relevant return information + results returned from the thread
17 int pthread_join(pthread_t thread, void **status);
18
```



1.1 PThread Attributes

Attributes:

- stack size
- inheritance
- joinable
- scheduling policy
- priority
- system/process scope

The attributes have default values, if we pass in NULL for attr argument in pthread_create(), then the thread will be created with default values for its attributes.

PThread Attribute Interface

```
1 int pthread_attr_init(pthread_attr_t *attr);
2 int pthread_attr_destroy(pthread_attr_t *attr);
3 // set or read a certain attribute value
4 pthread_attr_t {set/get}{attribute};
```



- By default, when a child thread is created, it's a joinable thread. The parent can wait and join the child thread, hence reaping the child thread.

- If the parent thread exits before joining its child threads, the child thread can become a zombie thread. After the child thread exits, it won't be reaped by its parent, so the memory assigned to it will not be freed.
- Detached thread
 - once detached, the thread cannot be joined
 - if a thread is detached, it will become equivalent to its parent thread
 - detached child thread can go on to execute after parent exit

```
1 // to detach a thread
2 pthread_detach();
```



To create a thread as detached thread

```
1 pthread_attr_t attr;
2
3 // ...
4
5 pthread_create(&tid, &attr, func, arg);
```



For a thread to exit execution:

```
1 pthread_exit(NULL);
```



Example to create detached thread:

```
#include <stdio.h>
#include <pthread.h>

void *foo (void *arg) { /* thread main */
    printf("Foobar!\n");
    pthread_exit(NULL);
}

int main (void) {
    int i;
    pthread_t tid;

    pthread_attr_t attr;
    pthread_attr_init(&attr); /* required!!! */
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    pthread_create(&tid, &attr, foo, NULL);

    return 0;
}
```

- Note that the PTHREAD_SCOPE_SYSTEM means the newly created thread will share resources equally with all other threads in the system.
The create statement should be:

```
1 pthread_create(&tid, &attr, foo, NULL);
```



1.2 Compile Pthreads

1. include header:

```
1 #include <pthread.h>
```



2. Compile source with -lpthread or -pthread

```
1 gcc -o main main.c -lpthread
```

```
2 gcc -o main main.c -pthread
```




3. Check return values of common functions

1.3 More Example

```
#include <stdio.h>
#include <pthread.h>
#define NUM_THREADS 4

void *hello (void *arg) { /* thread main */
    printf("Hello Thread\n");
    return 0;
}

int main (void) {
    int i;
    pthread_t tid[NUM_THREADS];
    for (i = 0; i < NUM_THREADS; i++) { /* create/fork threads */
        pthread_create(&tid[i], NULL, hello, NULL);
    }
    for (i = 0; i < NUM_THREADS; i++) { /* wait/join threads */
        pthread_join(tid[i], NULL);
    }
    return 0;
}
```



```

#include <stdio.h>
#include <pthread.h>
#define NUM_THREADS 4

void *threadFunc(void *pArg) { /* thread main */
    int *p = (int*)pArg;
    int myNum = *p;
    printf("Thread number %d\n", myNum);
    return 0;
}

int main(void) {
    int i;
    pthread_t tid[NUM_THREADS];
    for(i = 0; i < NUM_THREADS; i++) { /* create/fork threads */
        pthread_create(&tid[i], NULL, threadFunc, &i);
    }
    for(i = 0; i < NUM_THREADS; i++) { /* wait/join threads */
        pthread_join(tid[i], NULL);
    }
    return 0;
}

```

Private Variable

Note that in this program pArg is passed as a pointer to i, myNum can be the value that i is when the child thread access it. It might have changed say from 1 to 2.

To correct this behavior, copy the i value each thread is supposed to take into a static array, then pass in pointers to array elements corresponding to each thread.

```

#define NUM_THREADS 4

void *threadFunc(void *pArg) { /* thread main */
    int myNum = *((int*)pArg);
    printf("Thread number %d\n", myNum);
    return 0;
}

int main(void) {
    int tNum[NUM_THREADS];
    // ...
    for(i = 0; i < NUM_THREADS; i++) { /* create/fork threads */
        tNum[i] = i;
        pthread_create(&tid[i], NULL, threadFunc, &tNum[i]);
    }
    // ...
}

```

YouTube video player

2. PThread Mutex

Pthread Mutex Syntax

- 1 pthread_mutex_t aMutex; // mutex type
- 2 // lock operation
- 3 // explicit lock
- 4 int pthread_mutex_lock(pthread_mutex_t *mutex);
- 5 // explicit unlock
- 6 int pthread_mutex_unlock(pthread_mutex_t *mutex);



Birrell:

Pthreads

```
list<int> my_list;
Mutex m;
void safe_insert(int i) {
    Lock(m) {
        my_list.insert(i);
    } // unlock;
}
```

```
list<int> my_list;
pthread_mutex_t m;
void safe_insert(int i) {
    pthread_mutex_lock(m);
    my_list.insert(i);
    pthread_mutex_unlock(m);
}
```

Other mutex operations

```
1  int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);
2  // mutex attributes specifies mutex behavior when a mutex is shared among processes
3  // the attr can be NULL, meaning follow the default behavior
4
5  // try to acquire the lock, if lock is free, then acquire it, if not it'll not block but return and
6  // notify the calling thread that the mutex is not available
7  int pthread_mutex_trylock(pthread_mutex_t *mutex);
8
9  // free mutex data structure
10 int pthread_mutex_destroy(pthread_mutex_t *mutex);
```



Mutex Safety Tips

- shared data should always be accessed through a single mutex!
- mutex scope must be visible to all!
- globally order locks
 - for all threads, lock mutexes in order
- always unlock a mutex
 - always unlock the correct mutex

3. Condition Variables

Basic PThread condition variable API:

```
1 pthread_cond_t aCond; // type of cond variable
2 int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);
3 int pthread_cond_signal(pthread_cond_t *cond);
4 int pthread_cond_broadcast(pthread_cond_t *cond);
5
```

Other condition variable operations:

```
1 // initiate a condition variable with attr attributes.
2 int pthread_cond_init(pthread_cond_t *cond, const pthread_condattr_t *attr);
3 int pthread_cond_destroy(pthread_cond_t *cond);
```

Safety Tips:

- do not forget to notify waiting threads!
 - predicate change => signal/broadcast correct condition variable
- When in doubt broadcast
 - but **performance loss**
- You do not need a mutex to signal / broadcast

PThread Condition Variable Example — Producer Consumer Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

#define BUF_SIZE 3      /* size of shared buffer */

int buffer[BUF_SIZE];  /* shared buffer */
int add = 0;           /* place to add next element */
int rem = 0;           /* place to remove next element */
int num = 0;           /* number elements in buffer */

pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER; /* mutex lock for buffer */
pthread_cond_t c_cons = PTHREAD_COND_INITIALIZER; /* consumer waits on cv */
pthread_cond_t c_prod = PTHREAD_COND_INITIALIZER; /* producer waits on cv */

void *producer (void *param);
void *consumer (void *param);
```

```

int main(int argc, char *argv[]) {

    pthread_t tid1, tid2; /* thread identifiers */
    int i;

    if (pthread_create(&tid1, NULL, producer, NULL) != 0) {
        fprintf(stderr, "Unable to create producer thread\n");
        exit (1);
    }

    if (pthread_create(&tid2, NULL, consumer, NULL) != 0) {
        fprintf(stderr, "Unable to create consumer thread\n");
        exit (1);
    }

    pthread_join(tid1, NULL); /* wait for producer to exit */
    pthread_join(tid2, NULL); /* wait for consumer to exit */
    printf ("Parent quitting\n");
}

```

```

void *producer (void *param) {
    int i;
    for (i = 1; i <= 20; i++) {
        pthread_mutex_lock (&m);
        if (num > BUF_SIZE) { /* overflow */
            exit(1);
        }
        while (num == BUF_SIZE) { /* block if buffer is full */
            pthread_cond_wait (&c_prod, &m);
        }
        buffer[add] = i; /* buffer not full, so add element */
        add = (add+1) % BUF_SIZE;
        num++;
        pthread_mutex_unlock (&m);

        pthread_cond_signal (&c_cons);
        printf ("producer: inserted %d\n", i); fflush (stdout);
    }

    printf ("producer quitting\n"); fflush (stdout);
    return 0;
}

```

```
void *consumer (void *param) {  
    int i;  
    while (1) {  
        pthread_mutex_lock (&m);  
        if (num < 0) { /* underflow */  
            exit (1);  
        }  
        while (num == 0) { /* block if buffer empty */  
            pthread_cond_wait (&c_cons, &m);  
        }  
        i = buffer[rem]; /* buffer not empty, so remove element */  
        rem = (rem+1) % BUF_SIZE;  
        num--;  
        pthread_mutex_unlock (&m);  
  
        pthread_cond_signal (&c_prod);  
        printf ("Consume value %d\n", i); fflush(stdout);  
    }  
}
```

Note here the mutex is first released then signal condition variables to avoid spurious wakeups.